


CSIE5123: Numerical Method

Vinsong

April 11, 2026

Abstract

The lecture note of 2026 Spring Numerical Method by professor 林智仁.



Contents

1	Floating-point systems	2
1.1	Floating-point basics	2
1.2	Guard digits	4
1.3	Cancelation	8
1.4	IEEE 754 Standard	10
1.5	Trap Handlers	16
1.6	Real Examples	21
2	Direct Methods for Linear Systems	23
2.1	LU Factorization	23
2.2	Cholesky Factorization	30
2.3	Condition of Matrices	33
3	BLAS	37
3.1	BLAS basic	37
3.2	Optimized BLAS	41

Chapter 1

Floating-point systems

Lecture 1

1.1 Floating-point basics

2025.12.09

This chapter is mainly based on the science of floating-point arithmetics which is based on the IEEE standard 754.

Remark. The floating-point number system here introduction is to help us redesign a better floating-point system while we are doing deep learning implementation.

1.1.1 Why learning floating-point operations?

Example. A one-variable problem

$$\min_x f(x) \quad \text{where } x \geq 0$$

In the normal program, we should set an **upper bound** of x , or x may be wrongly increased to ∞ . We have to find the largest representable number in the computer. Or how to define the “infinity” in the computer?

Example. A ten-variable problem

$$\min_{\mathbf{x}} f(\mathbf{x}) \quad \text{where } x_i \geq 0, i = 1, 2, \dots, 10$$

We want to know how many are zeros, we may use

```
1 for (int i = 0; i < 10; i++)
2     if (x[i] == 0) count++;
3     // count the number of zeros
```

But people say that don't do the comparison of floating-point due to the precision issue.

```
1 double epsilon = 1.0e-12;
2 for (int i = 0; i < 10; i++)
3     if (x[i] <= epsilon) count++;
```

Which is better? How to chose “epsilon”? Can't do the comparison of floating-point? We need to understand the floating-point representation.

1.1.2 Floating-point Format

We know `float` (single precision): 4 bytes and `double` (double precision) 8 bytes in C/C++.

Note. We usually spare the bits for deep learning system to save memory and speed up the training process.

Definition 1.1.1 (format). A floating-point system requires

- a base β
- precision p
- significand (mantissa) $d.d \dots d$

Example.

$$\begin{aligned} 0.1 &= 1.00 \times 10^{-1} & (\beta = 10, p = 3) \\ &\approx 1.1001 \times 2^{-4} & (\beta = 2, p = 5) \end{aligned}$$

exponent: -1 and -4 ; significand: 1.00 and 1.1001 .

We let $e_{\max}$ to be the largest exponent and $e_{\min}$ to be the smallest exponent. Then, β^p is the possible significand, and $e_{\max} - e_{\min} + 1$ is the possible exponent.

$$\lceil \log_2(e_{\max} - e_{\min} + 1) \rceil + \lceil \log_2(\beta^p) \rceil + 1$$

bits for storing a floating-point number. 1 bit for the sign.

In the practical it is more complicated.

Definition 1.1.2 (Normalized floating-point number). A normalized floating-point number is a number whose significand is in the form of $1.d_1d_2 \dots d_{p-1}$.

Example. 0.1 is not; 1.00×2^{-3} is

Now the issue will become **represent of zero**. We naturally use $1.0 \times \beta^{e_{\min}-1}$ to represent zero, but it is not allowed in computer system.

Remark. We usually reserve some special numbers to represent 0 , ∞ , and NaN.

1.1.3 Relative error and Ulp

Example. When $\beta = 10, p = 3$, for 3.14159

We represent it as 3.14×10^0 , and the error is $|3.14159 - 3.14| = 0.00159 = 0.159 \times 10^{-2}$ i.e.

- 10^{-2} is the unit of the last place.
- ulps: 0.159 is the unit in the last place

The absolute error is 0.159 ulps.

Definition 1.1.3 (ulps). The unit in the last place (ulps) is the unit value of the least significant digit in a floating-point number.

We also use “relative error” to measure the error. The relative error is $0.00159/3.14159 \approx 0.0005$. For a number $d.d \dots d \times \beta^e$, the largest error cause by rounding is

$$0.\underbrace{0 \dots 0}_{p-1} \beta' \times \beta^e \quad \beta' = \beta/2$$

So, Error = $\frac{\beta}{2} \times \beta^{-p} \times \beta^e$ and

$$1 \times \beta^e \leq \text{original number} < \beta \times \beta^e$$

The relative error will be between

$$\frac{\frac{\beta}{2} \times \beta^{-p} \times \beta^e}{\beta^e} \text{ and } \frac{\frac{\beta}{2} \times \beta^{-p} \times \beta^e}{\beta^{e+1}}$$

So,

$$\text{relative error} \leq \frac{\beta}{2} \times \beta^{-p} \quad (1.1.1)$$

Definition 1.1.4 (machine epsilon). The machine epsilon is the upper bound of the relative error of a floating-point number $\epsilon = \frac{\beta}{2} \times \beta^{-p} = \beta^{1-p}/2$

Example. $x = 12.35 \Rightarrow \tilde{x} = 1.24 \times 10^1$ ($p = 3, \beta = 10$)

- 1 ulps = $0.01 \times 10^1 = 0.1$, $\epsilon = 1/2 \cdot 10^{-2} = 0.005$
- error = $0.05 = 0.005 \times 10^1 = 0.5$ ulps
- relative error = $0.05/12.35 \approx 0.004 = 0.8\epsilon$
- $8x = 98.8$, $8\tilde{x} = 9.92 \times 10^1$ error = 4.0 ulps, relative error = $0.04 = 0.8\epsilon$

Note. ulps and ϵ is used interchangeably usually.

1.2 Guard digits

In some situation, the order of floating-point computation and rounding may not have a huge effect on the final result. So, usually we choose to

$$\boxed{\text{Round}} \Rightarrow \boxed{\text{Compute}}$$

to let the computation to be more efficient. But in some cases, the error may be huge.

Example.

$$\begin{aligned} x &= 10.1 \\ y &= 9.93 \\ x - y &= 0.17 \end{aligned}$$

- Round and then compute:

$$10.1 - 9.93 = 1.01 \times 10^1 - 0.99 \times 10^1 = 0.02 \times 10^1 = 2.00 \times 10^{-1}$$

- error = $2.00 \times 10^{-1} - 0.17 = 0.03$
- ulps = $0.01 \times 10^{-1} = 10^{-3}$
- error = 30 ulps
- relative error

$$= \frac{0.03}{0.17} = \frac{3}{17}$$

- Compute and then round:

$$10.1 - 9.93 = 0.17 \approx 1.7 \times 10^{-1}$$

- error = $1.7 \times 10^{-1} - 0.17 = 0$

So, we have to know how big will the error be if we choose to round first then compute.

Theorem 1.2.1. Using p precision with base β for $x - y$, the relative error can be as large as $\beta - 1$

Proof. Let

$$x = 1.0 \dots 0, \quad y = 0.\underbrace{\eta \dots \eta}_{p \text{ digits}}, \quad \eta = \beta - 1$$

The correct solution is $x - y = \beta^{-p}$. Under some rounding scheme, we have the solution

$$1.0 \dots 0 - 0.\underbrace{\eta \dots \eta}_{p-1 \text{ digits}} = \beta^{-p+1}$$

Note. Above is a bad rounding just ignore the last η

Relative error is

$$\frac{|\beta^{-p+1} - \beta^{-p}|}{\beta^{-p}} = \beta - 1$$

■

Example. ($p = 3, \beta = 10$) $x = 1.00, y = 0.999, x - y = 0.001 = 10^{-3}$

The compute solution is $1.00 \times 10^0 - 0.99 \times 10^0 = 0.01 \times 10^0 = 0.01$. The relative error is

$$\frac{|0.01 - 0.001|}{0.001} = 9$$

Such a huge error occurs when the two numbers are close to each other. So, we need to use **guard digits** to store the intermediate result to avoid such a huge error.

Definition 1.2.1 (guard digits). Guard digits are extra digits used in the intermediate steps of a calculation to reduce the effect of rounding errors and improve the accuracy of the final result.

In here, the guard digits is **p increased by 1 in the device for addition and subtraction**

For the above example, we have the computation become

$$1.010 \times 10^1 - 0.993 \times 10^1 = 0.017 \times 10^1 = 0.17$$

note that $0.17 = 1.70 \times 10^{-1}$ can be stored in the system with $p = 3$ and $\beta = 10$. The relative error becomes 0.

Note. We only allocate more digit precision for the calculation device (i.e. the subtraction process) to store the intermediate result.

Example. ($p = 3, \beta = 10$) $x = 110, y = 8.59, x - y = 101.41$

$$\begin{aligned} 110 - 8.59 &= 1.100 \times 10^2 - 0.085 \times 10^2 \\ &= 1.015 \times 10^2 \approx 1.01 \times 10^2 \end{aligned}$$

Relative error is

$$\frac{|1.01 \times 10^2 - 101.41|}{101.41} \approx 0.003$$

$$\epsilon = \frac{1}{2} \cdot 10^{-2} = 0.005.$$

Theorem 1.2.2. Using $p + 1$ digits for $x - y \Rightarrow$ relative rounding error $< 2\epsilon$ (ϵ : machine epsilon)

Proof. We first make some assumption

- $x > y$
- $x = x_0 . x_1 x_2 \dots x_{p-1} \times \beta^0$, the proof will be similar for not β^0

For the first two condition,

1. If $y = y_0 . y_1 y_2 \dots y_{p-1}$ no error
2. If $y = 0.y_1 y_2 \dots y_p \Rightarrow$ 1 guard digit, exact $x - y$ rounded to a closest number $\Rightarrow$ error $< \epsilon$

In general, $y = 0.0 \dots 0 y_{k+1} \dots y_{k+p}$, we let $\bar{y}$: truncate y to $p + 1$ digits.

$$|y - \bar{y}| < (\beta - 1)(\beta^{-p-1} + \beta^{-p-2} + \dots + \beta^{-p-k}) \quad (1.2.1)$$

After the truncation, we have to compute

$$x - \bar{y}$$

Then we rounded the result get

$$x - \bar{y} + \delta$$

Thus

$$\text{error} \leq 0. \underbrace{0 \dots 0}_{p-1 \text{ digits}} (\beta/2) \quad (\text{worst case})$$

Therefore,

$$|\delta| \leq (\beta/2) \cdot \beta^{-p} \quad (1.2.2)$$

Then the error is

$$(x - y) - (x - \bar{y} + \delta) = \bar{y} - y - \delta$$

Now, we consider three cases:

1° Case 1: $x - y \geq 1$, from (1.2.1) and (1.2.2), we have the relative error

$$\begin{aligned}
 \frac{|\bar{y} - y - \delta|}{x - y} &\leq \frac{|\bar{y} - y - \delta|}{1} \\
 &\leq \beta^{-p}[(\beta - 1)(\beta^{-1} + \beta^{-2} + \dots + \beta^{-k}) + \beta/2] \\
 &= \beta^{-p}[(\beta - 1)\beta^{-k}(1 + \dots + \beta^{k-1}) + \beta/2] \\
 &= \beta^{-p}[(\beta - 1)\beta^{-k} \cdot \frac{\beta^k - 1}{\beta - 1} + \beta/2] \\
 &= \beta^{-p}[(1 - \beta^{-k}) + \beta/2] \\
 &< \beta^{-p}[1 + \beta/2] \leq 2\epsilon
 \end{aligned}$$

2° Case 2: $x - \bar{y} \leq 1$, enough digits to store $x - \bar{y}$, so the $\delta = 0$, and the relative error is

$$\frac{|\bar{y} - y|}{x - y}$$

The smallest $x - y$ is

$$1.0 - 0.\underbrace{0\dots 0}_k \rho \dots \rho > (\beta - 1) \overbrace{(\beta^{-1} + \beta^{-2} + \dots + \beta^{-k})}^{1.0 - 0\dots 01}$$

we have k zeros, p ρ and $\rho = \beta - 1$, from (1.2.1), we have the relative error

$$\begin{aligned}
 &\leq \frac{|\bar{y} - y|}{(\beta - 1)(\beta^{-1} + \beta^{-2} + \dots + \beta^{-k})} \\
 &< \frac{(\beta - 1)\beta^{-p}(\beta^{-1} + \beta^{-2} + \dots + \beta^{-k})}{(\beta - 1)(\beta^{-1} + \beta^{-2} + \dots + \beta^{-k})} = \beta^{-p} = 2\epsilon
 \end{aligned}$$

3° Case 3: $x - y < 1$ but $x - \bar{y} > 1$

If $x - \bar{y} = 1.\underbrace{0\dots 1}_p \Rightarrow x - y \geq 1$: a contradiction.

Because

$$|y - \bar{y}| < \beta^{-p} = 0.\underbrace{0\dots 1}_p$$

then

$$x - y = \underbrace{(1 + \beta^{-p})}_{x - \bar{y}} - \underbrace{|y - \bar{y}|}_{< \beta^{-p}} > 1$$

Proof is completed. ■

Conclusion: we can add “some” (not only 1) guard digits to make the relative error small enough. Especially when the two numbers are close to each other. To add one bit on the adder is very cheap.

Lecture 2

1.3 Cancellation

2026.03.01

Cancellation is an important source of numerical instability. It occurs when we subtract two nearly equal numbers.

1.3.1 Catastrophic Cancellation

Example. $b = 3.34$, $a = 1.22$, $c = 2.28$, calculate $b^2 - 4ac$

- True value: $b^2 - 4ac = 0.0292$
- Calculated value: $b^2 \approx 11.2$, $4ac \approx 11.1$, so $b^2 - 4ac \approx 0.1$

The Error = $0.1 - 0.0292 = 0.0708$

- Computed answer = $0.1 = 1.00 \times 10^{-1}$
- ulps = $0.01 \times 10^{-1} = 10^{-3}$
- $0.0708 \approx 70$ ulps

The error is huge when we subtract two nearly equal numbers. This is called **catastrophic cancellation**. We can use **benign cancellation** to avoid this problem. For example, we use guard digits to get the relative error being small.

1.3.2 Benign Cancellation

However, guard digits can only reduce the error of operations between two **exact number**. In this case, b^2 and $4ac$ already contain errors. We can use **reformulation** to avoid this problem. For example,

Example.

$$\frac{-b - \sqrt{b^2 - 4ac}}{2a} \quad (1.3.1)$$

If $b^2 \gg 4ac$, there is no problem for $\sqrt{b^2 - 4ac} \approx |b|$, but $-b + \sqrt{b^2 - 4ac}$ will cause catastrophic cancellation if $b > 0$.

We can reformulate by multiplying $-b - \sqrt{b^2 - 4ac}$ if $b > 0$.

$$\frac{2c}{-b + \sqrt{b^2 - 4ac}} \quad (1.3.2)$$

Now we can use (1.3.1) if $b < 0$ and (1.3.2) if $b > 0$ to avoid catastrophic cancellation.

Example. Calculate $x^2 - y^2$.

Assume $x \approx y$, $(x - y)(x + y)$ is much more accurate than $x^2 - y^2$ because x^2 and y^2 maybe rounded, so $x^2 - y^2$ may cause catastrophic cancellation. But, $x - y$ can be calculated by guard digits.

Remark. It is difficult to avoid all catastrophic cancellation but possible to remove some. A catastrophic cancellation is replaced by a benign cancellation.

Example. Calculate

$$A = \sqrt{s(s-a)(s-b)(s-c)}, \quad s = \frac{a+b+c}{2} \quad (1.3.3)$$

a, b, c are the lengths of the triangle.

If $a \approx b + c$,

$$s = \frac{a+b+c}{2} \approx a$$

and $s - a$ may cause catastrophic cancelation. A new formula is by Kahan[1986]: for $a \geq b \geq c$,

$$A = \frac{1}{4} \sqrt{(a + (b + c))(c - (a - b))(c + (a - b))(a + (b - c))} \quad (1.3.4)$$

Remark. Reformulation is a powerful tool to avoid catastrophic cancelation. It only works if guard digit is used.

1.3.3 Exactly Rounded Operations

Usually we choose to

$$\boxed{\text{Round}} \Rightarrow \boxed{\text{Compute}}$$

However, it may be not accurate. So, we introduce **exactly rounded** operations, which is to compute exactly then rounded to the nearest.

Definition (definition of rounding). There are two types of rounding:

Definition 1.3.1 (Rounding up). When the result's last digit $\geq$ half of the base, we round up, otherwise, we round down.

Definition 1.3.2 (Rounding even). Round to the closest value with even least significant digit.

Example. Round up

- $11.5 \Rightarrow 12$
- $12.5 \Rightarrow 13$

Round even

- $11.5 \Rightarrow 12$
- $12.5 \Rightarrow 12$

Theorem 1.3.1 (Reiser and Knuth). Let

$$x_0 = x, \quad x_1 = (x_0 \ominus y) \oplus y, \quad \dots, \quad x_n = (x_{n-1} \ominus y) \oplus y$$

If $\oplus, \ominus$ is an exactly rounded using even rounding, then

$$x_n = x \quad \forall n \text{ or } x_n = x_1 \quad \forall n \geq 1$$

Example. $\beta = 10, p = 3, x = 1.00, y = -0.555$

$$x - y = 1.555, x \ominus y = 1.56, (x \ominus y) + y = 1.56 - 0.555 = 1.005$$

- Rounding up:

$$x_1 = (x \ominus y) \oplus y = \textcolor{red}{1.01}, \quad x_1 - y = 1.565 \Rightarrow 1.57, \quad (x_1 \ominus y) + y = 1.57 - 0.555 = 1.015$$

$$x_2 = (x_1 \ominus y) \oplus y = \textcolor{red}{1.02}$$

Increase until $x_n = 9.45$.

- Rounding even:

$$x_1 = (x \ominus y) \oplus y = \textcolor{red}{1.00}, \quad x_1 - y = 1.555 \Rightarrow 1.56, \quad (x_1 \ominus y) + y = 1.56 - 0.555 = 1.005$$

$$x_2 = (x_1 \ominus y) \oplus y = \textcolor{red}{1.00}$$

Note. For the implementation, an array of words or floating-points used too much memory. Goldberg [1990] proposed a method to use only **two guard digits and one sticky bit**, the result is exact same as exactly rounded operations.

1.4 IEEE 754 Standard

1.4.1 Basics

IEEE 754 is the most used standard for floating-point arithmetic in computers nowadays. There are two of floating-point standard

- 754: specify $\beta = 2, p = 24$ for single precision, $\beta = 2, p = 53$ for double precision.
- 854: $\beta = 2$ or 10

Note. In standard 854, it accept 2, 10 because

- $\beta = 2$ smaller β causes smaller relative error.
- $\beta = 10$ is most used.

Consider

$$\beta = 16, p = 1 \text{ versus } \beta = 2, p = 4$$

which both used 4 bits to store the significand, but there ϵ is different

$$\epsilon_{16} = \frac{16}{2} \cdot 16^{-1} = \frac{1}{2}, \quad \epsilon_2 = \frac{2}{2} \cdot 2^{-4} = \frac{1}{16}$$

However, $\beta = 16$ is used by IBM/370 due to two reasons:

- 1° Assume 4 bytes = 32 bits, $\beta = 16, p = 6$
significand:

$$4 \times 6 = 24 \text{ bits}$$

exponent:

$$32 - 24 - 1 = 7 \text{ bits (1 bit for sign)}$$

The range of exponent is $16^{-2^6} \rightarrow 16^{2^6} = 2^{2^8}$.

If we instead use $\beta = 2$ and to keep the same range of exponent, then

9 bits ($-2^8 \rightarrow 2^8$) for exponent

and

$32 - 9 - 1 = 22$ bits for significand

Same exponents, less significand for $\beta = 2$ (24 vs. 22)

2° The cost of shifting: If $\beta = 16$, less frequently to adjust exponents when adding or subtracting two numbers. But nowadays the cost of shifting is not important since the computation speed.

1.4.2 IEEE standard: Significand and Exponent

For single precision, $\beta = 2, p = 24$ (23 bits as normalized), exponent 8 bits, and 1 bit for sign. Thus,

$32 = 1$ (sign bit) + 8 (exponent) + 23 (normalized significand)

Example. $176.625 = 1.0101100101 \times 2^7$

0 10000110 010110010100000000000000

1 of $1 \dots$ is not stored since the normalized.

Note. Here we use the bias of exponent

$10000110 = 128 + 4 + 2 = 134, 134 - 127 = 7$

Note that exponent may be negative, but here we don't use a sign bit for exponents

In this standard, it use rounding even

Binary	rounded	reason
10.00011	10.00	(< 1/2, down)
10.00110	10.01	(> 1/2, up)
10.11100	11.00	(1/2, up)
10.10100	10.10	(1/2, down)

Here is the summary of all standard:

IEEE	Fortran	C	Bits	Exp.	Mantissa
Single	REAL*4	float	32	8	24
Single-extended			44	≤ 11	32
Double	REAL*8	double	64	11	53
Double-extended	REAL*10	long double	≥ 80	≥ 15	≥ 64

Observe the single extended

$44 \neq 11$ (bits for exponent) + 32 (bits for significand)

Hardware implementation of extended precision normal don't use a hidden bit.

Minimal normalized positive number

$$1 \times 2^{-126} \approx 1.17 \times 10^{-38}$$

$$e_{\min} = -126$$

8 bits for exponent: 0 to 255. IEEE uses a biased approach for exponents

$$(0 \text{ to } 255) - 127 = -127 \text{ to } 128$$

There are some cases for exponent:

- -127: denormalized numbers, and 0
- -126 to 127: exponent for normalized numbers
- 128: special quantity (NaN, ∞)

Thus,

$$e_{\min} = -126, \quad e_{\max} = 127$$

Note. The reason for not using $e_{\min} = -127$ is that $1/2^{e_{\min}}$ not overflow, $1/2^{e_{\max}}$ underflow.

1.4.3 IEEE standard: Extended Precision

Some may use more digits for intermediate calculations, which is called **extended precision**. For example, binary-decimal conversion, or some calculators use 13 digits for intermediate calculations but round to 10 digits for display.

Another example is that writing 9 digits is enough for short. Though $10^8 > 2^{24}$, 8 digits are not enough. From Section 2.1.2 of Goldberg [1990], in reading the 9-digit number, if extended precision is available, an efficient algorithm can be done so that the original binary representation is recovered.

1.4.4 IEEE standard: Exactly Rounded Operations

IEEE standard requires results of addition, subtraction, multiplication and division exactly rounded. The reason of doing so is to leave the portability of software.

Also, square root, remainder, conversion between integer and floating-point, internal formats and decimal are exactly rounded.

IEEE does not require transcendental functions to be exactly rounded (e.g. $\sin, \cos, \exp$). They cannot specify the precision because they are arbitrarily long

1.4.5 IEEE standard: Special Quantities

On some computer, (e.g. IBM 370) there is no special quantity, the value like $\sqrt{-4} = 2$ and it print a error message. But IEEE standard has special quantities

NaN (Not a Number)

which means not every pattern of bits is a valid number. Here is the summary of all cases:

Exponent	significand	represents
$e = e_{\min} - 1$	$f = 0$	$+0, -0$
$e = e_{\min} - 1$	$f \neq 0$	$0.f \times 2^{e_{\min}}$
$e_{\min} \leq e \leq e_{\max}$		$1.f \times 2^e$
$e = e_{\max} + 1$	$f = 0$	$\pm\infty$
$e = e_{\max} + 1$	$f \neq 0$	NaN

Sometimes even 0/0 occurs, the program can continue, which is better than stop and print an error message.

Example. Finding $f(x) = 0$, try different x 's, even $0/0$ happens, other values may be ok.

If $b^2 - 4ac < 0$,

$$\frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$$

return NaN, and $-b \pm \sqrt{b^2 - 4ac}$ return NaN as well.

Remark. In general when a NaN is in an operation, result is NaN

Here are some examples of producing NaN:

Operation	NaN by
+	$\infty + (-\infty)$
$\times$	$0 \times \infty$
/	$0/0, \infty/\infty$
REM	$x \text{ REM } 0, \infty \text{ REM } y$
$\sqrt{}$	$\sqrt{x}$ when $x < 0$

1.4.6 IEEE standard: Infinity

$\beta = 10, p = 3, e_{\max} = 98, x = 3 \times 10^{70}, x^2$ overflow and replaced by ∞ . Note that

$$0/0 = \text{NaN}, 1/0 = \infty, -1/0 = -\infty \Rightarrow \text{nonzero divided by 0 is } \infty \text{ or } -\infty$$

In the computation, we replace ∞ with x , let $x \rightarrow \infty$.

Example.

$$\frac{x}{x^2 + 1} \text{ vs } \frac{1}{x + x^{-1}}$$

For $\frac{x}{x^2 + 1}$, if x is large, x^2 overflows, so $\frac{x}{\infty} = 0$ but not $\frac{\frac{1}{x}}{\text{not } 0}$. For $\frac{1}{x + x^{-1}}$, if x is large, $\frac{1}{x}$ is ok, which is better.

When for $x = 0$,

$$\frac{1}{x + x^{-1}} = \frac{1}{0 + \infty} = \frac{1}{\infty} = 0$$

is also better.

Note. If no infinity arithmetic, an extra instruction is needed to test if $x = 0$. This may interrupt the pipeline.

1.4.7 IEEE standard: Signed Zero

To discuss why we need signed zero, first, it is available with 1 sign bit. If no signed zero, $1/(1/x)$ will fail when $x = \pm\infty$

$$\begin{aligned} x = \infty, 1/x = 0, 1/0 &= +\infty \\ x = -\infty, 1/x = 0, 1/0 &= +\infty \end{aligned}$$

Definition 1.4.1 (Comparison of signed zero). In IEEE standard, $+0 = -0$.

± 0 is useful for some case

Example. We can define

$$\log x \equiv \begin{cases} -\infty & x = 0 \\ \text{NaN} & x < 0 \end{cases}$$

Definition 1.4.2 (Underflow). Underflow means a value smaller than the smallest floating-point number occurs.

When face a small underflow negative number, $\log x$ should return NaN. However, in the definition above,

$$\boxed{x \text{ underflow negative}} \Rightarrow \boxed{x \stackrel{\text{round}}{=} 0} \Rightarrow \boxed{\log x = -\infty}$$

With signed zero, we can define $\log x$ as

$$\log x \equiv \begin{cases} -\infty & x = +0 \\ \text{NaN} & x = -0 \\ \text{NaN} & x < 0 \end{cases}$$

Positive underflow will round to $+0$, and negative underflow will round to -0 .

Example (Complex arithmetic).

$$\frac{1}{\sqrt{z}} \quad \text{and} \quad \sqrt{\frac{1}{z}}$$

If $z = -1$,

$$\frac{1}{\sqrt{-1}} = \frac{1}{i} = -i \quad \sqrt{\frac{1}{-1}} = i$$

Thus, $1/\sqrt{z} \neq \sqrt{1/z}$. Which happen because

$$i^2 = (-i)^2 = -1$$

However, by some restrictions (or ways of calculation), they can be equal.

If $z = -1 = -1 + 0i$,

$$\frac{1}{-1 + 0i} = \frac{-1 - 0i}{(-1)^2 + (0)^2} = -1 - 0i$$

then

$$\sqrt{\frac{1}{z}} = \sqrt{-1 - 0i} = -i$$

with signed zero, we can let $\sqrt{1/z} = 1/\sqrt{z}$.

But there is a disadvantage of signed zero, if $x = +0$ and $y = -0$, then

$$\frac{1}{x} = \infty, \quad \frac{1}{y} = -\infty$$

1.4.8 IEEE standard: Denormalized Numbers

Example. Consider

$$\beta = 10, p = 3, e_{\min} = -98, x = 6.87 \times 10^{-97}, y = 6.81 \times 10^{-97}$$

x, y are ok but $x - y$ underflow and round to 0 even though $x \neq y$.

Note. It is important to preserve the property of

$$x = y \Leftrightarrow x - y = 0$$

with this case: if $(x \neq y) \{ z = 1/(x-y); \}$ The statement is true, but z becomes ∞ .

IEEE use denormalized numbers to solve this problem, which is the most controversial part of IEEE standard. If denormalized number is used, 0.6×10^{-98} is also a floating-point number.

We do not store the leading 1 for $1.\dots$. Recall the valid range of exponent, $e \geq e_{\min}$, and we have $1.d\dots d \times 2^e$. For denormalized numbers, we allow $e = e_{\min} - 1$, and the corresponding value is

$$0.d\dots d \times 2^{e_{\min}}$$

Remark. The reason why not using

$$1.d\dots d \times 2^{e_{\min}-1}$$

is that then we cannot represent

$$0.0d\dots d \times 2^{e_{\min}}$$

which is even smaller than the smallest normalized number.

Example. $6.87 \times 10^{-97} - 6.81 \times 10^{-97} = 0.06 \times 10^{-97}$

is underflow due to the cancelation.

Here is an example of using denormalized numbers:

Example.

$$\frac{a+bi}{c+di} = \frac{(a+bi)(c-di)}{(c+di)(c-di)} = \frac{ac+bd}{c^2+d^2} + \frac{bc-ad}{c^2+d^2}i$$

If c or $d > \sqrt{\beta}\beta^{e_{\max}/2}$, there will be overflow.

Definition 1.4.3 (Overflow). Overflow means a value larger than the largest floating-point number occurs.

Lemma 1.4.1 (Smith's lemma).

$$\frac{a+bi}{c+di} = \begin{cases} \frac{a+b(d/c)}{c+d(d/c)} + \frac{b-a(d/c)}{c+d(d/c)}i & \text{if } (|d| < |c|) \\ \frac{b+a(c/d)}{d+c(c/d)} + \frac{-a+b(c/d)}{d+c(c/d)}i & \text{if } (|d| \geq |c|) \end{cases}$$

This is a reformulation to avoid overflow. However, using Smith's formula, some issues may occur without denormalized numbers:

If

$$a = 2 \times 10^{-98}, \quad b = 1 \times 10^{-98}, \quad c = 4 \times 10^{-98}, \quad d = 2 \times 10^{-98}$$

then

$$\begin{aligned} \frac{d}{c} &= 0.5, & c + d \cdot \frac{d}{c} &= 5 \times 10^{-98} \\ b \cdot \frac{d}{c} &= 1 \times 10^{-98} \times 0.5 = \mathbf{0}, & a + b \cdot \frac{d}{c} &= 2 \times 10^{-98} \end{aligned}$$

Solution = 0.4, which is wrong.

If denormalized numbers are used, 0.5×10^{-98} can be stored

$$a + b \cdot \frac{d}{c} = 2 \times 10^{-98} + 0.5 \times 10^{-98} = 2.5 \times 10^{-98}$$

Solution = 0.5, which is correct.

Note. Usually hardware does not support denormalized numbers directly, we use software to simulate it.

Note. Program may be slow when there is a lot of underflow.

1.5 Trap Handlers

1.5.1 Exception, Flags, Trap handlers

When we face with some special situations (e.g. divide by zero), we have to know what happens. For this purpose, IEEE requires vendors to provide a way to get status flags. **overflow**, **underflow**, **division by zero**, **invalid operation**, **inexact**.

Definition (IEEE exception). IEEE defines five exceptions:

Definition 1.5.1 (Overflow). Larger than the maximal floating-point number.

Definition 1.5.2 (Underflow). Smaller than the minimal floating-point number.

Definition 1.5.3 (Division by zero). Nonzero divided by zero.

Definition 1.5.4 (Invalid operation). $\infty + (-\infty)$, $0 \times \infty$, $0/0$, ∞/∞ , $x \text{ REM } 0$, $\infty \text{ REM } y$, $\sqrt{x}$, $x < 0$, any comparison involves a NaN.

Remark. Invalid returns NaN; But NaN may not be from invalid operations.

Definition 1.5.5 (Inexact). The result is not exact: $\beta = 10, p = 3, 3.5 \times 4.2 = 14.7$ exact, $3.5 \times 4.3 = 15.05 \Rightarrow 15.0$ not exact.

Remark. Inexact exception is raised so often, usually we ignore it.

Here is the summary of all exceptions and their corresponding handlers:

Exception	when trap disabled	argument to handler
overflow	$\pm\infty$ or $\pm 1.1 \dots 1 \times 2^{e_{\max}}$	$\text{round}(x \times 2^{-\alpha})$
underflow	$0, \pm 2^{e_{\min}}$, or denormalized	$\text{round}(x \times 2^{\alpha})$
division by zero	∞	operands
invalid	NaN	operands
inexact	$\text{round}(x)$	$\text{round}(x)$

Definition 1.5.6 (Trap handler). Special subroutines to handle exception, which can be designed by users.

In the above table, “when trap disabled” means results of operations if trap handlers not used.

Example. $\alpha = 192$, for single precision, $\alpha = 1536$ for double precision. We can’t store x since it is too large/small.

1.5.2 Compiler Operations

Compiler may provide a way so the program stops if an exception occurs. For example, SUN’s C compiler provides a command

```
1 -ftrap=t
```

t: %all, %none, common, [%no%]invalid, [%no%]overflow, [%no%]underflow, [%no%]division, [%no%]inexact

Note. common: invalid, division by zero, and overflow

when compile you will see

```
1 Note: IEEE floating-point exception flags raised:
2   Inexact; Underflow;
3 See the Numerical Computation Guide, ieee_flags(3M)
```

Note. gcc doesn’t have the functionality to explicitly detect exceptions. It only provides

- `-fno-trapping-math`: The default if `-ftrapping-math`. Setting this option may allow faster code if one relies on “non-stop” IEEE arithmetic.
- `-ftrapv`: Generates traps for signed overflow on addition, subtraction, multiplication.

1.5.3 Trap Handlers

For example,

```
1 do {
2   ...
3 } while { not x >= 100; }
```

If $x = \text{NaN}$, an infinite loop will occur. Any comparison involving NaN is wrong. Thus $x \geq 100$ returns false. A trap handler can be installed to abort it.

Another example is that calculate $x_1 \times \dots \times x_n$ may overflow in the middle, but fine in the end.

```
1 for (i = 1; i <= n; i++)
2     p = p * x[i];
```

If trap handlers are not used, the result will be ∞ and wrong, i.e.

$$x_1 \times \dots \times x_r, r \leq n \text{ overflow, } x_1 \times \dots \times x_n \text{ in the range}$$

Note. There might be a solution,

$$e^{\sum \log(x_i)}$$

but quite unaccurate and cost more.

A trap handler can be installed to save the intermediate result

```
1 for (i = 1; i <= n; i++) {
2     if (p * x[i] overflow) { // trap handler to catch overflow
3         p = p * pow(10, -a);
4         count = count + 1 ;
5     }
6     p = p * x[i];
7 }
8 p = p * pow(10, a*count) ;
```

Lecture 3

On SUN machines, there is an additional math library: `libsunmath.a`. In this library, there are `exp2`, `exp10`, ..., `ieee_flags`, `ieee_handler`, and `ieee_retrospective`. 2026.03.07

```
1 #include <stdio.h>
2 #include <sys/ieee.h>
3 #include <sunmath.h>
4 #include <siginfo.h>
5 #include <ucontext.h>
6
7 void handler(int sig, siginfo_t *sip, ucontext_t *uap)
8 {
9     unsigned code, addr;
10
11     code = sip->si_code;
12     addr = (unsigned) sip->si_addr;
13     fprintf(stderr, "fp exception %x at address %x \n", code, addr);
14 }
15 int main()
16 {
17     double x;
18
19     /* trap on common floating point exceptions */
20     if (ieee_handler("set", "common", handler) != 0)
```

```

21     printf("Did not set exception handler \n");
22
23     /* cause an underflow exception (not reported) */
24     x = min_normal();
25     printf("min_normal = %g \n", x);
26     x = x / 13.0;
27     printf("min_normal / 13.0 = %g \n", x);
28
29     /* cause an overflow exception (reported) */
30     x = max_normal();
31     printf("max_normal = %g \n", x);
32     x = x * x;
33     printf("max_normal * max_normal = %g \n", x);
34
35     ieee_retrospective(stderr);
36     return 0;
37 }

```

The results of the above code are as follows:

```

min_normal = 2.22507e-308
min_normal / 13.0 = 1.7116e-309
max_normal = 1.79769e+308
fp exception 4 at address 10d0c
max_normal * max_normal = 1.79769e+308
Note: IEEE floating-point exception flags raised:
    Inexact; Underflow;
IEEE floating-point exception traps enabled:
    overflow; division by zero; invalid operation;
See the Numerical Computation Guide, ieee_flags(3M), ieee_handler(3M)

```

In this example, we can see that the underflow exception is not reported, due to the denormalized numbers in the IEEE 754 standard. In this handler, we simply print the something, but in practice, we can do more error handling in this kind of subroutine.

Example. Calculate x^n , n : integer.

```

1  double pow(double x, int n)
2  {
3      double tmp = x, ret = 1.0;
4      for(int t=n; t>0; t/=2)
5      {
6          if (t % 2 == 1) ret *= tmp;
7          tmp = tmp * tmp;
8      }
9      return ret;
10 }

```

When $n > 0$, the above code can do

$$x^n = \prod_{i=0}^{\lfloor \log_2 n \rfloor} x^{2^i \cdot (n \bmod 2^{i+1} > 0)}$$

to calculate x^n in $O(\log n)$ time. But if $n < 0$, we need to use

$$x^n = \left(\frac{1}{x}\right)^{-n} = \frac{1}{x^{-n}}$$

Due to there is already error on $1/x$, $\text{pow}(1/x, -n)$ is less accurate than $1/\text{pow}(x, -n)$.

However, there is another problem: if $\text{pow}(x, -n)$ cause an underflow, but $1/\text{pow}(x, -n)$ does not overflow

$$e_{\min} = -126, \quad 2^{-e_{\min}} = 2^{126} < 2^{127} = 2^{e_{\max}}$$

The solution is

1° turn off overflow & underflow trap enable bits, save overflow & underflow status bits.

2° Calculate $\text{pow}(x, -n)$.

- If neither overflow nor underflow status is set $\Rightarrow$ restore them.
- If one is set, restore & calculate $\text{pow}(1/x, -n)$, which causes correct exception to occur restore them.

But practically, the calculation of $\text{pow}()$ is more complicated.

Example. Calculate $\arccos$ using $\arctan$.

Theorem 1.5.1.

$$\arccos x = 2 \arctan \sqrt{\frac{1-x}{1+x}}$$

Proof.

$$\begin{aligned} \cos \theta = x &= 2 \cos^2 \frac{\theta}{2} - 1 = 1 - 2 \sin^2 \frac{\theta}{2} \\ \Rightarrow \cos \frac{\theta}{2} &= \sqrt{\frac{1+x}{2}}, \quad \sin \frac{\theta}{2} = \sqrt{\frac{1-x}{2}} \\ \Rightarrow \tan \frac{\theta}{2} &= \sqrt{\frac{1-x}{1+x}} \Rightarrow \theta = 2 \arctan \sqrt{\frac{1-x}{1+x}} \end{aligned}$$

■

Consider $x = -1$

$$\arctan(\infty) = \frac{\pi}{2} \Rightarrow \arccos(-1) = \pi$$

There is a small problem $(1-x)/(1+x)$ cause the divide-by-zero flag set though $\arccos(-1)$ is not exceptional. The solution is to save divide-by-zero flag, restore it after $\arccos$ computation.

1.6 Real Examples

1.6.1 Same Code but Different Architectures

```

1  #include <stdio.h>
2
3  int main()
4  {
5      float a = 123.123;
6      printf("%.10f\n", a);
7      printf("%.10f\n", a*a);
8
9      a = 123.125;
10     printf("%.10f\n", a);
11     printf("%.10f\n", a*a);
12 }

```

```

$ gcc test.c; ./a.out
123.1230010986
15159.2734375000
123.1250000000
15159.7656250000
$ gcc -m32 test.c; ./a.out
123.1230010986
15159.2733995339
123.1250000000
15159.7656250000

```

Note. `-m 32` generates code for a 32-bit environment.

Same code gives different results under 32 and 64-bit environments.

Note. On 32 bit, 387 floating-point coprocessor is used. From gcc manual, “The temporary results are computed in 80-bit precision instead of the precision specified by the type, resulting in slightly different results compared to most of other chips.” In other words, they somehow **violate IEEE standard**.

The number 123.123 has infinite digits after transformed to binary, assigning the compiler option can make the result consistent.

Example. `-mfpmath=387` to let the 64-bit machine run like a 32-bit one.

```

$ gcc -mfpmath=387 test.c; ./a.out
123.1230010986
15159.2733995339
123.1250000000
15159.7656250000

```

Example. `-ffloat-store` make the 32-bit machine like a 64-bit one Manual of the option, “Do not store floating-point variables in registers, and inhibit other options that might change whether a floating-point value is taken from a register or memory.”

```
$ gcc -ffloat-store test.c; ./a.out
123.1230010986
15159.2734375000
123.1250000000
15159.7656250000
$ gcc -ffloat-store -m32 test.c; ./a.out
123.1230010986
15159.2734375000
123.1250000000
15159.7656250000
```

1.6.2 Order of Operations

For the same code, other issues such as order of operations can also affect results.

Example. Consider running a real example using a machine learning software [LIBSVM](#).

In the O0 optimization level:

```
$ g++ -O0 svm-train.c svm.cpp -o svm-train -lm
$ ./svm-train -c 100 -e 0.00001 heart_scale
.....*..*
optimization finished, #iter = 2872
nu = 0.148045
obj = -2526.925470, rho = 1.145512
nSV = 107, nBSV = 9
Total nSV = 107
```

In the Ofast optimization level:

```
$ g++ -Ofast svm-train.c svm.cpp -o svm-train -lm
$ ./svm-train -c 100 -e 0.00001 heart_scale
.....*..*
optimization finished, #iter = 2910
nu = 0.148045
obj = -2526.925470, rho = 1.145510
nSV = 107, nBSV = 9
Total nSV = 107
```

Some compiler optimizations may change the order of operations to improve performance. On default settings for 64-bit environments, O0 to O3 produce the same results. But from gcc manual, `-Ofast` “disregards strict standards compliance”.

`-mfpmath=387` is even more sensitive to optimizations. To produce the same results with `-mfpmath=387`, we need to disable all optimizations due to more complicated interactions with registers and memory.

Chapter 2

Direct Methods for Linear Systems

2.1 LU Factorization

2.1.1 Solving Dense Linear Systems

Solving a linear system $Ax = b$ is an important problem in numerical method. In this kind of problem, we assume A is a dense matrix, i.e. most

$$A_{ij} \neq 0$$

and we allocate space to store

$$A_{ij}, \quad \forall i, \forall j$$

2.1.2 Triangular Systems

Start from a triangular system,

$$\begin{pmatrix} l_{11} & \\ l_{21} & l_{22} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \end{pmatrix}$$

if $l_{11}, l_{22} \neq 0$,

$$x_1 = \frac{b_1}{l_{11}}, \quad x_2 = \frac{b_2 - l_{21}x_1}{l_{22}}$$

Theorem 2.1.1. In general,

$$x_i = \frac{\left(b_i - \sum_{j=1}^{i-1} l_{ij}x_j \right)}{l_{ii}}$$

This is called “**forward substitution**”, which requires $O(n^2)$ operations. Similarly, we can solve an upper triangular system using “**backward substitution**”.

Theorem 2.1.2. In general,

$$x_i = \frac{\left(b_i - \sum_{j=i+1}^n u_{ij}x_j \right)}{u_{ii}}$$

2.1.3 LU Factorization

Example. Consider the following system

$$\begin{aligned} 3x_1 + 5x_2 &= 9 \\ 6x_1 + 7x_2 &= 4 \end{aligned}$$

By Gaussian elimination,

$$\begin{aligned} 3x_1 + 5x_2 &= 9 \\ -3x_2 &= -14 \end{aligned}$$

Actually, the process is in the below

$$\begin{aligned} \begin{pmatrix} 3 & 5 \\ 6 & 7 \end{pmatrix} &= \begin{pmatrix} 1 & 0 \\ 2 & 1 \end{pmatrix} \begin{pmatrix} 3 & 5 \\ 0 & -3 \end{pmatrix} \\ \Rightarrow \begin{pmatrix} 1 & 0 \\ 2 & 1 \end{pmatrix} \begin{pmatrix} y_1 \\ y_2 \end{pmatrix} &= \begin{pmatrix} 9 \\ 4 \end{pmatrix} \Rightarrow \begin{pmatrix} y_1 \\ y_2 \end{pmatrix} = \begin{pmatrix} 9 \\ -14 \end{pmatrix} \end{aligned}$$

Then solve

$$\begin{pmatrix} 3 & 5 \\ 0 & -3 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = \begin{pmatrix} 9 \\ -14 \end{pmatrix}$$

is same as solving the original system.

Example. Consider

$$A = \begin{pmatrix} 1 & 4 & 7 \\ 2 & 5 & 8 \\ 3 & 6 & 10 \end{pmatrix}$$

We do the Gaussian Transformation to get the upper triangular matrix, if

$$M_1 = \begin{pmatrix} 1 & 0 & 0 \\ -2 & 1 & 0 \\ -3 & 0 & 1 \end{pmatrix}$$

then

$$M_1 A = \begin{pmatrix} 1 & 4 & 7 \\ 0 & -3 & -6 \\ 0 & -6 & -11 \end{pmatrix}$$

Note that

$$\begin{pmatrix} -3 & -6 \\ -6 & -11 \end{pmatrix} = \begin{pmatrix} 5 & 8 \\ 6 & 10 \end{pmatrix} - \begin{pmatrix} 2 \\ 3 \end{pmatrix} \begin{pmatrix} 4 & 7 \end{pmatrix}$$

Likewise,

$$\begin{aligned} M_2 &= \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & -2 & 1 \end{pmatrix} \\ \Rightarrow M_2(M_1 A) &= \begin{pmatrix} 1 & 4 & 7 \\ 0 & -3 & -6 \\ 0 & 0 & 1 \end{pmatrix} \end{aligned}$$

Hence, the whole Gaussian elimination process is

$$U = M_{n-1} \cdots M_2 M_1 A$$

$$A = (M_1^{-1}(M_2^{-1} \cdots (M_{n-1}^{-1}U))) = LU$$

M_i^{-1} is a lower triangular matrix,

$$M_i = \begin{pmatrix} 1 & 0 & 0 \\ \vdots & \ddots & 0 \\ m_{i1} & \cdots & 1 \end{pmatrix}, \quad M_i^{-1} = \begin{pmatrix} 1 & 0 & 0 \\ \vdots & \ddots & 0 \\ -m_{i1} & \cdots & 1 \end{pmatrix}$$

$$M_1 = \begin{pmatrix} 1 & 0 & 0 \\ -2 & 1 & 0 \\ -3 & 0 & 1 \end{pmatrix} \Rightarrow M_1^{-1} = \begin{pmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ 3 & 0 & 1 \end{pmatrix}$$

$$M_2 = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & -2 & 1 \end{pmatrix} \Rightarrow M_2^{-1} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 2 & 1 \end{pmatrix}$$

$$M_1^{-1}M_2^{-1} = \begin{pmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ 3 & 0 & 1 \end{pmatrix} \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 2 & 1 \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ 3 & 2 & 1 \end{pmatrix}$$

In the implementation, we do not store $M_1, \dots, M_{n-1}$, they can be stored in the lower triangular part of A .

The reason of choosing LU factorization than Gaussian elimination is that Gaussian elimination overwrites the original matrix. If you do not overwrite A , need another array (the same size as A and $L + U$)

Gaussian elimination update right-hand side together. To handle **multiple right-hand sides**, LU is the most effective way. The algorithm is as follows,

```

1  for k=1:n-1
2      A(k+1:n,k) = A(k+1:n,k)/A(k,k);
3      A(k+1:n, k+1:n) = A(k+1:n, k+1:n) - A(k+1:n,k)*A(k,k+1:n);
4  end

```

The number of operations is

$$\approx \sum_{k=1}^{n-1} 2k^2 \approx 2 \cdot \frac{n(n-1)(2n-1)}{6} = \frac{2}{3}n^3$$

Note. The constant 2 is due to one multiplication and one subtraction.

2.1.4 Problem of LU Factorization

If matrix A has a zero diagonal such as

$$A = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$$

then we can't generate M_1 . Even if we can do LU factorization, numerical error is sometimes large.

Example. Assume $\beta = 10, p = 3$, **no guard digit**. Solving

$$\begin{pmatrix} 0.001 & 1.00 \\ 1.00 & 2.00 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = \begin{pmatrix} 1.00 \\ 3.00 \end{pmatrix}$$

After LU factorization, we have

$$L = \begin{pmatrix} 1 & 0 \\ 1000 & 1 \end{pmatrix}, \quad U = \begin{pmatrix} 0.001 & 1 \\ 0 & -1000 \end{pmatrix}$$

$2 - 1000 = -998 \approx -1000$ in the floating-point arithmetic. However, the real solution is

$$\begin{pmatrix} 0.001 & 1 \\ 0 & -998 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = \begin{pmatrix} 1 \\ -997 \end{pmatrix}$$

$$x_2 = \frac{997}{998}$$

$$x_1 = 1000(1 - \frac{997}{998}) = \frac{1000}{998}$$

But using our system, we have

$$\begin{pmatrix} 1 & 0 \\ 1000 & 1 \end{pmatrix} \begin{pmatrix} y_1 \\ y_2 \end{pmatrix} = \begin{pmatrix} 1.00 \\ 3.00 \end{pmatrix}$$

$$y_1 = 1.00, \quad y_2 = 3.00 - 1000 \approx -1000$$

$$\begin{pmatrix} 0.001 & 1 \\ 0 & -1000 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = \begin{pmatrix} 1.00 \\ -1000 \end{pmatrix} \tag{2.1.1}$$

Then we get

$$x_1 = 0, \quad x_2 = 1$$

$(0, 1)$ is very different from the real solution $(1000/998, 997/998)$.

Remark. A small pivot may cause large numerical errors.

In the eqs. (2.1.1), if x_2 contains an errors, because $1 \gg 0.001$, x_1 is seriously affected. Hence, if we interchange the two rows of A

$$\begin{pmatrix} 0.001 & 1 \\ 1 & 2 \end{pmatrix} \Rightarrow \begin{pmatrix} 1 & 2 \\ 0.001 & 1 \end{pmatrix}$$

then

$$L = \begin{pmatrix} 1 & 0 \\ 0.001 & 1 \end{pmatrix}, \quad U = \begin{pmatrix} 1 & 2 \\ 0 & 1 \end{pmatrix}$$

$$\begin{pmatrix} 1 & 0 \\ 0.001 & 1 \end{pmatrix} \begin{pmatrix} y_1 \\ y_2 \end{pmatrix} = \begin{pmatrix} 3.00 \\ 1.00 \end{pmatrix}$$

$$y_1 = 3.00, \quad y_2 = 1.00$$

therefore, we solve

$$\begin{pmatrix} 1 & 2 \\ 0 & 1 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = \begin{pmatrix} 3.00 \\ 1.00 \end{pmatrix}$$

$$x_1 = 1.00, \quad x_2 = 1.00$$

which is much closer to the real solution.

Remark. We can use permutation matrix P to represent the row interchanges.

Example.

$$P = \begin{pmatrix} 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 \end{pmatrix}$$

This means we do the exchange of

$$4 \rightarrow 1, \quad 1 \rightarrow 2, \quad 3 \rightarrow 3, \quad 2 \rightarrow 4$$

2.1.5 Pivoting Strategies

Pivoting in LU:

Example.

$$A = \begin{pmatrix} 3 & 17 & 10 \\ 2 & 4 & -2 \\ 6 & 18 & -12 \end{pmatrix}$$

We first find the **largest absolute value** in the first column: 6, so we do

$$P_1 = \begin{pmatrix} 0 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \end{pmatrix}$$

then

$$P_1 A = \begin{pmatrix} 6 & 18 & -12 \\ 2 & 4 & -2 \\ 3 & 17 & 10 \end{pmatrix} \Rightarrow M_1 = \begin{pmatrix} 1 & 0 & 0 \\ -1/3 & 1 & 0 \\ -1/2 & 0 & 1 \end{pmatrix}$$

$$M_1 P_1 A = \begin{pmatrix} 6 & 18 & -12 \\ 0 & -2 & 2 \\ 0 & 8 & 4 \end{pmatrix}$$

$$|8| > |-2|$$

$$P_2 = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 1 & 0 \end{pmatrix}$$

then

$$M_2 = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 1/4 & 1 \end{pmatrix}$$

$$M_2 P_2 M_1 P_1 A = \begin{pmatrix} 6 & 18 & -12 \\ 0 & 8 & 16 \\ 0 & 0 & 6 \end{pmatrix}$$

The general form is

$$M_{n-1} P_{n-1} \cdots M_2 P_2 M_1 P_1 A = U$$

Now we have

$$Ax = b$$

$$\Rightarrow (M_{n-1} P_{n-1} \cdots M_2 P_2 M_1 P_1)^{-1} Ax = (M_{n-1} P_{n-1} \cdots M_2 P_2 M_1 P_1)^{-1} b$$

hence, we can solve

$$Ux = (M_{n-1} P_{n-1} \cdots M_2 P_2 M_1 P_1)^{-1} b$$

2.1.6 Pivoting Strategies: Implementation

However storing M_i and P_i is not efficient, we can try to overwrite A to store M_i and P_i .

$$\begin{pmatrix} 3 & 17 & 10 \\ 2 & 4 & -2 \\ 6 & 18 & -12 \end{pmatrix} \xrightarrow{p(1)=3} \begin{pmatrix} 6 & 18 & -12 \\ 2 & 4 & -2 \\ 3 & 17 & 10 \end{pmatrix}$$

then

$$\begin{pmatrix} 6 & 18 & -12 \\ \boxed{1/3} & -2 & 2 \\ \boxed{1/2} & 8 & 16 \end{pmatrix} \xrightarrow{p(2)=3} \begin{pmatrix} 6 & 18 & -12 \\ 1/2 & 8 & 16 \\ 1/3 & -2 & 2 \end{pmatrix}$$

Remark. For the boxed elements, we usually fill 0 in the upper triangular part, but to store the multiplier if that row more efficient. We can do it by filling the multiplier in the lower triangular part (dividing it by the pivot in its column).

In the above procedure we use an array p to store the permutation. Then we try to swap two rows of A together.

$$\begin{pmatrix} 0 & 0 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \end{pmatrix} \begin{pmatrix} 3 & 17 & 10 \\ 2 & 4 & -2 \\ 6 & 18 & -12 \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 \\ 1/2 & 1 & 0 \\ 1/3 & -1/4 & 1 \end{pmatrix} \begin{pmatrix} 6 & 18 & -12 \\ 0 & 8 & 16 \\ 0 & 0 & 6 \end{pmatrix}$$

Hence, it is possible to get the general form

$$PA = LU$$

where P is the permutation matrix. Now we can overwrite A to store L and U , and use an array p to store the permutation. To get x , we have

$$PAx = Pb$$

then we solve

$$(LU)x = Pb$$

We can get the general form of the algorithm as follows,

```

1  for k=1:n-1
2      find m
3      swap A(k,1:n) and A(m, 1:n)
4      p(k) = m;
5      A(k+1:n,k) = A(k+1:n,k)/A(k,k);
6      A(k+1:n, k+1:n) = A(k+1:n, k+1:n) - A(k+1:n,k)*A(k,k+1:n);
7  end

```

Theorem 2.1.3.

$$PA = LU$$

and L is the lower-triangular matrix obtained in the algorithm.

Proof. By the definition of P , we hope to get

$$\begin{aligned} PA &= P_{n-1} \cdots P_2 P_1 A \\ &= P_{n-1} \cdots P_1^{-1} M_1^{-1} \cdots P_{n-1}^{-1} M_{n-1}^{-1} U \\ &= LU \end{aligned}$$

Consider i -th column of

$$P_{n-1} \cdots P_1^{-1} M_1^{-1} \cdots P_{n-1}^{-1} M_{n-1}^{-1}$$

Note that for

$$P_j^{-1}, M_j^{-1}, \quad j > i$$

the i -th column is always

$$\begin{pmatrix} \vdots \\ 0 \\ \cdots & 1 & \cdots \\ 0 \\ \vdots \end{pmatrix}$$

Thus,

$$(P_{n-1} \cdots P_1 P_1^{-1} M_1^{-1} \cdots P_{n-1}^{-1} M_{n-1}^{-1})_{:,i} = (P_{n-1} \cdots P_1^{-1} M_1^{-1} \cdots P_i^{-1} M_i^{-1})_{:,i}$$

Note.

$$(M_i^{-1})_{:,i} = \begin{pmatrix} 0 \\ \vdots \\ 0 \\ 1 \\ m_{i+1,i} \\ \vdots \\ m_{n,i} \end{pmatrix}$$

So, by similar argument, we have

$$(P_{n-1} \cdots P_1 P_1^{-1} M_1^{-1} \cdots P_i^{-1} M_i^{-1})_{:,i} = (P_{n-1} \cdots P_{i+1} M_i^{-1})_{:,i} \quad (2.1.2)$$

Here is the detail

$$(P_i^{-1} M_i^{-1}) = \begin{pmatrix} \vdots \\ 0 \\ x' \\ x' \\ \vdots \\ x' \end{pmatrix}$$

because P_i only swap two rows in $i, \dots, n$ -th rows, then

$$(M_{i-1}^{-1} P_i^{-1} M_i^{-1})_{:,i} = \begin{pmatrix} 1 & & & \\ & x & 1 & \\ \cdots & \vdots & 0 & 1 \\ & \vdots & \vdots & \ddots \\ & x & 0 & \end{pmatrix} \begin{pmatrix} \vdots \\ 0 \\ x' \\ x' \\ \vdots \\ x' \end{pmatrix} = \begin{pmatrix} \vdots \\ 0 \\ x' \\ x' \\ \vdots \\ x' \end{pmatrix}$$

which does not change the i -th column.

Therefore, we have

$$\begin{aligned}(P_{n-1} \cdots P_1 P_1^{-1} M_1^{-1} \cdots P_i^{-1} M_i^{-1})_{:,i} &= (P_{n-1} \cdots P_1 P_1^{-1} P_2^{-1} \cdots P_i^{-1} M_i^{-1})_{:,i} \\ &= (P_{n-1} \cdots P_{i+1} M_i^{-1})_{:,i}\end{aligned}$$

where eliminate $P_i^{-1} = P_i$. ■

Note. Eqs. (2.1.2) is exactly get in the algorithm, and the $P_{n-1}, \dots, P_{i+1}$ swapping them into the corresponding position.

2.2 Cholesky Factorization

There are some special cases of matrix A . For example, diagonal, tri-diagonal, banded matrices, positive definite matrices, etc.

Example. A banded matrix,

$$\begin{pmatrix} 1 & 2 & 3 & 0 & 0 & 0 \\ 4 & 5 & 6 & 7 & 0 & 0 \\ 8 & 9 & 1 & 2 & 3 & 0 \\ 0 & 4 & 5 & 6 & 7 & 8 \\ 0 & 0 & 9 & 1 & 2 & 3 \\ 0 & 0 & 0 & 4 & 5 & 6 \end{pmatrix}$$

Many numerical methods are designed for special matrices such as **symmetric positive definite matrices**.

Definition 2.2.1 (positive definite). An $n \times n$ matrix A is positive definite if

$$x^T A x > 0, \quad \forall x \in \mathbb{R}^n, x \neq 0$$

Theorem 2.2.1. A is symmetric positive definite if and only if all A 's eigenvalues are positive.

Theorem 2.2.2 (symmetric positive definite). A is SPD $\Rightarrow$ all diagonal elements are positive, all principle sub-matrices are positive definite.

Definition 2.2.2 (principle sub-matrix). Let A be an $n \times n$ matrix. A principal submatrix is a submatrix obtained by deleting some rows and the same columns. Formally, given a subset of indices $I \subseteq \{1, 2, \dots, n\}$, the principal submatrix A_I is formed by selecting the rows and columns indexed by I .

Theorem 2.2.3 (Cholesky factorization). If A is SPD, then there exists a unique lower triangular matrix L such that

$$A = LL^T$$

The above theroem is called **Cholesky factorization**.

2.2.1 Cholesky Factorization

Theorem 2.2.4 (Cholesky factorization). We have

$$\begin{aligned} A &= \begin{pmatrix} \alpha & v^T \\ v & B \end{pmatrix} \quad (\alpha > 0) \\ &= \begin{pmatrix} \sqrt{\alpha} & 0 \\ \frac{v}{\sqrt{\alpha}} & I \end{pmatrix} \begin{pmatrix} \sqrt{\alpha} & \frac{v^T}{\sqrt{\alpha}} \\ 0 & B - \frac{vv^T}{\alpha} \end{pmatrix} \\ &= \begin{pmatrix} \sqrt{\alpha} & 0 \\ \frac{v}{\sqrt{\alpha}} & I \end{pmatrix} \begin{pmatrix} \sqrt{\alpha} & 0 \\ 0 & B - \frac{vv^T}{\alpha} \end{pmatrix} \begin{pmatrix} 1 & \frac{v^T}{\alpha} \\ 0 & I \end{pmatrix} \end{aligned}$$

If

$$B - \frac{vv^T}{\alpha} = \bar{L}\bar{L}^T$$

i.e. it is also SPD, then we have

$$\begin{aligned} \begin{pmatrix} \sqrt{\alpha} & 0 \\ \frac{v}{\sqrt{\alpha}} & I \end{pmatrix} \begin{pmatrix} \sqrt{\alpha} & 0 \\ 0 & \bar{L}\bar{L}^T \end{pmatrix} \begin{pmatrix} 1 & \frac{v^T}{\alpha} \\ 0 & I \end{pmatrix} &= \left(\begin{pmatrix} \sqrt{\alpha} & 0 \\ \frac{v}{\sqrt{\alpha}} & I \end{pmatrix} \begin{pmatrix} 1 & 0 \\ 0 & \bar{L} \end{pmatrix} \right) \left(\begin{pmatrix} 1 & 0 \\ 0 & \bar{L}^T \end{pmatrix} \begin{pmatrix} 1 & \frac{v^T}{\alpha} \\ 0 & I \end{pmatrix} \right) \\ &= \begin{pmatrix} \sqrt{\alpha} & 0 \\ \frac{v}{\sqrt{\alpha}} & \bar{L} \end{pmatrix} \begin{pmatrix} \sqrt{\alpha} & \frac{v^T}{\sqrt{\alpha}} \\ 0 & \bar{L}^T \end{pmatrix} \end{aligned}$$

Note. If we can prove that $B - \frac{vv^T}{\alpha}$ is also SPD, then the same procedure can be applied to $B - \frac{vv^T}{\alpha}$, and we can get the Cholesky factorization of A .

Lemma 2.2.1. If A is SPD, then $B - \frac{vv^T}{\alpha}$ is also SPD.

Proof. For the symmetricity, it is obvious. For the positive definiteness, we have for any $x \neq 0$,

$$\begin{aligned} x^T \left(B - \frac{vv^T}{\alpha} \right) x &= x^T Bx - \frac{(v^T x)^2}{\alpha} \\ &= \begin{pmatrix} -\frac{v^T x}{\sqrt{\alpha}} & x^T \end{pmatrix} \begin{pmatrix} 0 \\ -\frac{v^T x}{\alpha} v + Bx \end{pmatrix} \\ &= \underbrace{\begin{pmatrix} -\frac{v^T x}{\sqrt{\alpha}} & x^T \end{pmatrix}}_{\mathbf{v}} \underbrace{\begin{pmatrix} \alpha & v^T \\ v & B \end{pmatrix}}_A \underbrace{\begin{pmatrix} -\frac{v^T x}{\sqrt{\alpha}} \\ x \end{pmatrix}}_{\mathbf{v}^T} > 0 \quad (\text{because } A \text{ is SPD}) \end{aligned}$$

■

Now, we can go through the implementation in the **outer product form**.

```

1 for k=1:n
2     A(k,k) = sqrt(A(k,k))
3     A(k+1:n,k) = A(k+1:n,k)/A(k,k)
4     for j=k+1:n
5         A(j:n,j) = A(j:n,j) - A(j:n,k)A(j,k)
6     end
7 end

```

The inner loop comes from

$$B_{:,j} - (v/\alpha)_j \cdot \frac{v}{\sqrt{\alpha}}$$

There is another implementation way, treat

$$A = LL^T$$

as a system of equations with unknowns L and solve it.

$$\begin{aligned} A_{ij} &= \sum_{k=1}^n L_{ik}(L^T)_{kj} = \sum_{k=1}^n L_{ik}L_{jk} = \sum_{k=1}^n L_{jk}L_{ik} \\ \Rightarrow A_{:,j} &= \sum_{k=1}^n L_{jk}L_{:,k} \end{aligned}$$

Because L is lower triangular, we have

$$L_{j,j+1} = \dots = L_{j,n} = 0$$

Thus,

$$\begin{aligned} A_{:,j} &= \sum_{k=1}^j L_{jk}L_{:,k} \\ L_{jj}L_{:,j} &= A_{:,j} - \sum_{k=1}^{j-1} L_{jk}L_{:,k} \end{aligned}$$

The $(j-1)$ -th column is already computed, so we can compute $L_{:,j}$ by let

$$L_{jj}L_{jj} = A_{j,j} - \sum_{k=1}^{j-1} L_{jk}L_{j,k}$$

Then

$$L_{j+1:n,j} = \frac{(A_{j+1:n,j} - \sum_{k=1}^{j-1} L_{jk}L_{j+1:n,k})}{L_{jj}}$$

In practical implementation we calculate

$$A_{j:n,j} - \sum_{k=1}^{j-1} L_{jk}L_{j:n,k}$$

first.

```

1  for j=1:n
2      v(j:n) = A(j:n,j)
3      for k=1:j-1
4          v(j:n) = v(j:n) - L(j,k)L(j:n,k)
5      end
6      L(j:n,j) = v(j:n)/sqrt(v(j))
7  end

```

Note.

$$v(j:n) = v(j:n) - L(j,k)L(j:n,k)$$

is a **apxy** operation.

For the time complexity,

$$\begin{aligned}
 \sum_{j=1}^n \sum_{k=1}^{j-1} 2(n-j+1) &= 2 \sum_{j=1}^n (j-1)(n-j+1) \\
 &= \sum_{j=1}^n (nj - n - j^2 + j + j - 1) \\
 &\approx 2 \left(n \cdot \frac{n(n+1)}{2} - \frac{n(n+1)(2n+1)}{6} \right) = O\left(\frac{n}{3}\right)
 \end{aligned}$$

Which is half of the time complexity of LU factorization.

Lecture 4

2.3 Condition of Matrices

2026.03.14

2.3.1 Vector and Matrix Norms

To calculate the error, we can separate into 3 types: scalar, vector, and matrix. For scalar, we can simply use

- absolute error:

$$|\tilde{\alpha} - \alpha|$$

- relative error:

$$\frac{|\tilde{\alpha} - \alpha|}{|\alpha|}$$

To evaluate the vector and matrix error, we need to define the vector and matrix norms, which is a kind of metric.

Definition (Vector Norm). Here are some common vector norms:

Definition 2.3.1 (l -norm).

$$\|x\|_l = \left(\sum_{i=1}^n |x_i|^l \right)^{\frac{1}{l}}$$

for instance,

Definition 2.3.2 (1-norm).

$$\|x\|_1 = \sum_{i=1}^n |x_i|$$

Definition 2.3.3 (∞ -norm).

$$\|x\|_\infty \equiv \lim_{l \rightarrow \infty} \|x\|_l$$

Theorem 2.3.1.

$$\|x\|_\infty = \max_{1 \leq i \leq n} |x_i|$$

Proof. We can induce by Squeeze Theorem. For any l , we have

$$\sqrt[l]{(\max |x_i|)^l} \leq \sqrt[l]{|x_1|^l + \dots + |x_n|^l} \leq \sqrt[l]{n(\max |x_i|)^l}$$

which implies

$$\max |x_i| \leq \|x\|_l \leq n^{\frac{1}{l}} \max |x_i|$$

By Squeeze Theorem ($l \rightarrow \infty$), we have

$$\|x\|_\infty = \max |x_i|$$

Proof is complete. ■

Example.

$$x = \begin{pmatrix} 1 \\ 1000 \\ 9 \end{pmatrix}, \quad \hat{x} = \begin{pmatrix} 1.1 \\ 99 \\ 11 \end{pmatrix}$$

$$\begin{aligned} \|\hat{x} - x\|_\infty &= 2, & \frac{\|\hat{x} - x\|_\infty}{\|x\|_\infty} &= 0.02, & \frac{\|\hat{x} - x\|_\infty}{\|\hat{x}\|_\infty} &= 0.0202 \\ \|\hat{x} - x\|_2 &= 2.238, & \frac{\|\hat{x} - x\|_2}{\|x\|_2} &= 0.0223, & \frac{\|\hat{x} - x\|_2}{\|\hat{x}\|_2} &= 0.0225 \end{aligned}$$

Corollary 2.3.1. For l -norm, all norms are equivalent to each other, which means there exist $c_1, c_2 > 0$ such that

$$c_1 \|x\|_{l_1} \leq \|x\|_{l_2} \leq c_2 \|x\|_{l_1}$$

for all $x \in \mathbb{R}^n$.

For instance,

$$\|x\|_2 \leq \|x\|_1 \leq \sqrt{n} \|x\|_2 \tag{2.3.1}$$

$$\|x\|_\infty \leq \|x\|_2 \leq \sqrt{n} \|x\|_\infty \tag{2.3.2}$$

$$\|x\|_\infty \leq \|x\|_1 \leq n \|x\|_\infty \tag{2.3.3}$$

Theorem 2.3.2 (norm). A norm should satisfy

$$\|A\| \geq 0$$

$$\|A + B\| \leq \|A\| + \|B\| \tag{2.3.4}$$

$$\|\alpha A\| = |\alpha| \|A\|$$

where α is a scalar.

Proof. Checking the middle one, we have

$$\|A + B\| = \max_{\|x\|=1} \|(A + B)x\| \leq \max_{\|x\|=1} \|Ax\| + \|Bx\| \leq \max_{\|x\|=1} \|Ax\| + \max_{\|x\|=1} \|Bx\|$$

The last step is because we can choose different x to maximize $\|Ax\|$ and $\|Bx\|$. ■

Hence, here we have the definition of matrix norm:

Definition 2.3.4 (Matrix Norm). A matrix norm is a function $\|\cdot\|$

$$\|A\|_l \equiv \max_{x \neq 0} \frac{\|Ax\|_l}{\|x\|_l} = \max_{\|x\|_l=1} \|Ax\|_l$$

From the definition, above we have

$$\frac{\|Ax\|}{\|x\|} \leq \|A\|, \quad \forall x$$

so

$$\|Ax\| \leq \|A\| \|x\|$$

Note. For the relative error,

$$\frac{|\hat{\alpha} - \alpha|}{|\alpha|}$$

but in practice, α is usually unknown, so we can only use

$$\frac{|\hat{\alpha} - \alpha|}{|\hat{\alpha}|}$$

If

$$\frac{|\hat{\alpha} - \alpha|}{|\hat{\alpha}|} \leq 0.1$$

then

$$\frac{1}{1.1} \frac{|\hat{\alpha} - \alpha|}{|\hat{\alpha}|} \leq \frac{|\hat{\alpha} - \alpha|}{|\alpha|} \leq \frac{1}{0.9} \frac{|\hat{\alpha} - \alpha|}{|\hat{\alpha}|}$$

By

$$|\alpha| - |\hat{\alpha}| \leq |\hat{\alpha} - \alpha| \leq 0.1|\hat{\alpha}|$$

$$|\alpha| \leq 1.1|\hat{\alpha}|$$

Similarly,

$$|\hat{\alpha}| - |\alpha| \leq 0.1|\hat{\alpha}|$$

$$|\alpha| \geq 0.9|\hat{\alpha}|$$

2.3.2 Condition of a Linear System

Solving a Linear System

$$\begin{pmatrix} 10 & 7 & 8 & 7 \\ 7 & 5 & 6 & 5 \\ 8 & 6 & 10 & 9 \\ 7 & 5 & 9 & 10 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{pmatrix} = \begin{pmatrix} 32 \\ 23 \\ 33 \\ 31 \end{pmatrix}, \text{ solution} = \begin{pmatrix} 1 \\ 1 \\ 1 \\ 1 \end{pmatrix}$$

Now we slightly modify the right-hand side to

$$\begin{pmatrix} 10 & 7 & 8 & 7 \\ 7 & 5 & 6 & 5 \\ 8 & 6 & 10 & 9 \\ 7 & 5 & 9 & 10 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{pmatrix} = \begin{pmatrix} 32.1 \\ 22.9 \\ 33.1 \\ 30.9 \end{pmatrix}, \text{ solution} = \begin{pmatrix} 9.2 \\ -12.6 \\ 4.5 \\ -1.1 \end{pmatrix}$$

We can see that **a small change can lead to a large error**. Or if we slightly modify the coefficient matrix

$$\begin{pmatrix} 10 & 7 & 8.1 & 7.2 \\ 7.08 & 5.04 & 6 & 5 \\ 8 & 5.98 & 9.89 & 9 \\ 6.99 & 4.99 & 9 & 9.98 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{pmatrix} = \begin{pmatrix} 32 \\ 23 \\ 33 \\ 31 \end{pmatrix}, \text{ solution} = \begin{pmatrix} -81 \\ 137 \\ -34 \\ 22 \end{pmatrix}$$

which is even worse.

Now, we consider the general case

$$Ax = b$$

1° First, we slightly modify right-hand side

$$A(x + \delta x) = b + \delta b$$

then we get

$$\delta x = A^{-1} \delta b \quad \Rightarrow \quad \|\delta x\| \leq \|A^{-1}\| \|\delta b\|$$

and we have

$$\|b\| \leq \|A\| \|x\|$$

so

$$\boxed{\frac{\|\delta x\|}{\|x\|} \leq \|A^{-1}\| \|A\| \frac{\|\delta b\|}{\|b\|}}$$

2° Next, we slightly modify the coefficient matrix

$$(A + \delta A)(x + \delta x) = b$$

$$Ax + A\delta x + \delta A(x + \delta x) = b$$

then we have

$$\delta x = -A^{-1} \delta A(x + \delta x)$$

imply the norm

$$\|\delta x\| \leq \|A^{-1}\| \|\delta A\| (\|x\| + \|\delta x\|)$$

so

$$\boxed{\frac{\|\delta x\|}{\|x + \delta x\|} \leq \|A\| \|A^{-1}\| \frac{\|\delta A\|}{\|A\|}}$$

Note. We here estimate only

$$\frac{\|\delta x\|}{\|x + \delta x\|} \text{ instead of } \frac{\|\delta x\|}{\|x\|};$$

since the discussion earlier.

Clearly, the error is related to $\|A\| \|A^{-1}\|$, which is called the **condition number** of A , a smaller condition number means a more stable system.

Chapter 3

BLAS

Lecture 5

3.1 BLAS basic

2026.03.21

3.1.1 BLAS: Basic Linear Algebra Subroutines

Usually we just use 90% of the time in 10% of the code. So, accelerating the 10% of the code can significantly improve the overall performance. In linear algebra, we use BLAS (Basic Linear Algebra Subroutines) to accelerate the performance of the code, which is a kind of “software” face of optimization.

Note. For GPU, TPU, and other hardware, those are the “hardware” face of optimization.

For example,

```
1 daxpy(n, alpha, p, inc, w, inc);
2 daxpy(n, malpha, q, inc, r, inc);
3
4 rtr = ddot(n, r, inc, r, inc);
5 rnorm = sqrt(rtr);
6 tnorm = sqrt(ddot(n, t, inc, t, inc));
```

- ddot: inner product
- daxpy: $ax + y$, x, y are vectors and a is a scalar

Note. Fortran is the first high level programming language, and it is still widely used in scientific computing. The BLAS library is come up before the C language.

3.1.2 Level 1 BLAS

The level 1 BLAS includes:

- dot product
- constant times a vector plus a vector
- rotation (don't need to know what this is now)
- copy vector x to vector y

- swap vector x and vector y
- length of a vector ($\sqrt{x_1^2 + \dots + x_n^2}$)
- sum of absolute values ($|x_1| + \dots + |x_n|$)
- constant times a vector
- index of element having maximum absolute value

These are all $O(n)$ operations.

```
1 dw = ddot(n, dx, incx, dy, incy)
```

$$w = \sum_{i=1}^n dx_{1+(i-1)incx} dy_{1+(i-1)incy}$$

```
1 dw = ddot(n, dx, 2, dy, 2)
```

$$w = dx_1 dy_1 + dx_3 dy_3 + \dots$$

Note. `sdot` is single precision, `ddot` is for double precision.

If we want $y = ax + y$

```
1 daxpy(n, da, dx, incx, dy, incy)
```

For researcher in this field, it is hard to decide which operations can be put into BLAS. Further, there is another problem: Originally, BLAS is designed for Fortran, but now we have many programming languages. For C calling Fortran

```
1 rtr = ddot_(&n, r, &inc, r, &inc);
```

- `ddot_`: calling Fortran subroutines (machine dependent)
- `&n`: call by reference for Fortran subroutine.
- Array index starts from 1 in Fortran, but starts from 0 in C. But this should not cause problems here.

There is a CBLAS interface for C, i.e. C call C library,

```
1 rtr = ddot(n, r, inc, r, inc);
```

We don't need to handle the issue.

Nowadays, we have BLAS libraries available on most computers, for a Linux system

```
$ ls /usr/lib| grep blas
libblas
libblas.a
libblas.so
libblas.so.3
libblas.so.3gf
libcblas.so.3
libcblas.so.3gf
libf77blas.so.3
libf77blas.so.3gf
```


- .a: static library
- .so: dynamic library
- .3: versus
- .3gf: g77 versus gfortran

3.1.3 Level 2 BLAS

Level 2 BLAS involves $O(n^2)$ operations, where n is the size of matrices.

Example. For a Matrix-vector product:

$$Ax : (Ax)_i = \sum_{j=1}^n A_{ij}x_j, \quad i = 1, \dots, m$$

This kind of operation using level 1 BLAS will be too slow

$$y = \alpha Ax + \beta y, \quad y = \alpha A^T x + \beta y, \quad y = \alpha \bar{A}^T x + \beta y$$

α, β are scalars, x, y are vectors, A is a matrix, and $\bar{A}^T$ is the conjugate transpose (or Hermitian transpose) of A . For real matrices,

$$\bar{A}^T = A^T$$

Lower- or upper-triangular matrix-vector products

$$x = Tx, x = T^T x, x = \bar{T}^T x$$

x is a vector and T is a lower or upper triangular matrix.

Example. Rank-one and rank-two updates

$$A = \alpha xy^T + A, H = \alpha x\bar{y}^T + \bar{\alpha} y\bar{x}^T + H$$

- H is a Hermitian matrix ($H = \bar{H}^T$, symmetric for real numbers)
- rank: # of independent rows (columns) of a matrix.

xy^T is a rank one matrix. For example,

$$\begin{pmatrix} 1 \\ 2 \end{pmatrix} \begin{pmatrix} 3 & 4 \end{pmatrix} = \begin{pmatrix} 3 & 4 \\ 6 & 8 \end{pmatrix}$$

This operation need $O(n^2)$ operations.

$xy^T + yx^T$ is a rank-two matrix

$$\begin{pmatrix} 1 \\ 2 \end{pmatrix} \begin{pmatrix} 3 & 4 \end{pmatrix} + \begin{pmatrix} 3 & 4 \\ 6 & 8 \end{pmatrix} = \begin{pmatrix} 3 & 4 \\ 6 & 8 \end{pmatrix}, \begin{pmatrix} 3 \\ 4 \end{pmatrix} \begin{pmatrix} 1 & 2 \end{pmatrix} + \begin{pmatrix} 3 & 6 \\ 4 & 8 \end{pmatrix} = \begin{pmatrix} 3 & 6 \\ 4 & 8 \end{pmatrix}$$

Example. Solution of triangular equations

$$x = T^{-1}y$$

$$\begin{pmatrix} T_{11} & & & \\ T_{21} & T_{22} & & \\ \vdots & \vdots & \ddots & \\ T_{n1} & T_{n2} & \cdots & T_{nn} \end{pmatrix} \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix} = \begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix}$$

The solution is

$$\begin{aligned} x_1 &= y_1/T_{11} \\ x_2 &= (y_2 - T_{21}x_1)/T_{22} \\ x_3 &= (y_3 - T_{31}x_1 - T_{32}x_2)/T_{33} \end{aligned}$$

Number of multiplications/divisions:

$$1 + 2 + \cdots + n = \frac{n(n+1)}{2}$$

3.1.4 Level 3 BLAS

Example. Matrix-matrix product

$$\begin{aligned} C &= \alpha AB + \beta C, & C &= \alpha A^T B + \beta C, \\ C &= \alpha AB^T + \beta C, & C &= \alpha A^T B^T + \beta C \end{aligned}$$

Example. Rank- k and rank- $2k$ updates.

Example. Multiplying a matrix by a triangular matrix

$$B = \alpha TB$$

Example. Solving triangular systems of equations with multiple right-hand side:

$$B = \alpha T^{-1}B$$

BLAS does not include subroutines for solving general linear systems or eigenvalues. Because these problems are not as common as the above operations.

3.2 Optimized BLAS

3.2.1 Memory Management

Here is a C program

```

1  #define n 3000
2  double a[n][n], b[n][n], c[n][n];
3
4  int main()
5  {
6      int i, j, k;
7      for (i = 0; i < n; i++)
8          for (j = 0; j < n; j++) {
9              a[i][j] = 1; b[i][j] = 1;
10             }
11
12     for (i = 0; i < n; i++)
13         for (j = 0; j < n; j++) {
14             c[i][j] = 0;
15             for (k = 0; k < n; k++)
16                 c[i][j] += a[i][k] * b[k][j];
17         }
18 }

```

```

$ gcc -O3 mat.c; time ./a.out
real    1m24.909s
user    1m24.534s
sys     0m0.193s

```

We now test in the MATLAB environment

```

$ matlab -singleCompThread
>> n = 3000;
>> A = randn(n,n); B = randn(n,n);
>> tic; C = A*B; toc
Elapsed time is 1.708523 seconds.

```

An issue about timing is elapsed time versus CPU time

```

>> A = randn(n,n); B = randn(n,n);
>> t = cputime; C = A*B; t = cputime - t
t = 1.3000

```

They are similar if no other jobs are running on this machine.

Results of using multi-threading (the default of MATLAB)

```


$ matlab
>> n = 3000;
>> A = randn(n,n); B = randn(n,n);
>> tic; C = A*B; toc

```

```
Elapsed time is 0.426942 seconds.
>> A = randn(n,n); B = randn(n,n);
>> t = cputime; C = A*B; t = cputime -t
t = 5.1200
```

The CPU time here is the total time used by all threads.

Optimized BLAS is an implementation that takes the advantage of memory hierarchies



Note. In the reality, the architecture of CPU is more complicated, and there are multiple levels of cache.

Definition 3.2.1 (Block algorithm). **Block algorithms** is a way to transfer sub-matrices between different levels of storage. They localize operations to achieve good performance.

Definition 3.2.2 (Page fault). **Page fault** is an event where an operand is not available in main memory and must be transported from secondary memory. Usually if a page is moved to the main memory, it overwrites page least recently used.

Example. $C = AB + C$, $n = 1,024$

Here is the assumption:

- a page 65,536 doubles = 64 columns
- 16 pages for each matrix
- 48 pages for three matrices
- available memory 16 pages
- matrices access: **column oriented**

Note.

$$A = \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix}$$

The access way of the elements:

column oriented: 1 3 2 4

- row oriented: 1 2 3 4

Access each row of A: 16 page faults, $1024/64 = 16$

Example. Here is approach 1:

```

1  for i =1:n
2      for j=1:n
3          for k=1:n
4              c(i,j) = a(i,k)*b(k,j)+c(i,j);
5          end
6      end
7  end

```

At each (i,j) : each row $a(i, 1:n)$ causes 16 page faults. Total:

$$1024^2 \times 16$$

at least **16 million page faults**.

Example. Here is approach 2:

```

1  for j=1:n
2      for k=1:n
3          for i=1:n
4              c(i,j) = a(i,k)*b(k,j)+c(i,j);
5          end
6      end
7  end

```

For each j , access all columns of A . A needs 16 pages, but B and C take spaces as well. So A must be read for every j . For each j , 16 page faults for A .

$$1024 \times 16$$

and C, B : 16 page faults. This approach has a dramatic improvement, which is **16,384 page faults**.

If we implement back to C,

```

1  #define n 3000
2  double a[n][n], b[n][n], c[n][n];
3
4  int main()
5  {
6      int i, j, k;
7      for (i = 0; i < n; i++)
8          for (j = 0; j < n; j++) {
9              a[i][j] = 1; b[i][j] = 1;
10             c[i][j] = 0;
11         }
12
13     for (j = 0; j < n; j++) {
14         for (k = 0; k < n; k++)

```

```

15     for (i = 0; i < n; i++)
16         c[i][j] += a[i][k]*b[k][j];
17     }
18 }

```

```

$ gcc -O3 mat1.c; time ./a.out
real    4m20.247s
user    4m19.761s
sys      0m0.154s

```

The performance is much worse than the previous one. The reason is that C language is row oriented.

Example. Here is approach 3 (block algorithm with $nb = 256$):

```

1  for j=1:nb:n
2      for k=1:nb:n
3          for jj=j:j+nb-1
4              for kk=k:k+nb-1
5                  c(:,jj) = a(:,kk)*b(kk,jj)+c(:,jj);
6              end
7          end
8      end
9  end

```

In MATLAB, $1:256:1025$ means 1, 257, 513, 769. Note that we calculate

$$\begin{pmatrix} A_{11} & \cdots & A_{14} \\ & \ddots & \\ A_{41} & \cdots & A_{44} \end{pmatrix} \begin{pmatrix} B_{11} & \cdots & B_{14} \\ & \ddots & \\ B_{41} & \cdots & B_{44} \end{pmatrix} = \begin{pmatrix} A_{11}B_{11} + \cdots + A_{14}B_{41} & \cdots \\ & \ddots & \\ & & \ddots \end{pmatrix}$$

We split the original matrix into 16 blocks, for each block 256×256

$$C_{11} = A_{11}B_{11} + \cdots + A_{14}B_{41}$$

$$C_{21} = A_{21}B_{11} + \cdots + A_{24}B_{41}$$

$$C_{31} = A_{31}B_{11} + \cdots + A_{34}B_{41}$$

$$C_{41} = A_{41}B_{11} + \cdots + A_{44}B_{41}$$

For each (j, k) , $B_{k,j}$ is used to add $A_{:,k}B_{k,j}$ to $C_{:,j}$. For instance:

$$C_{11} \leftarrow C_{11} + A_{11}B_{11}$$

$$\vdots$$

$$C_{41} \leftarrow C_{41} + A_{41}B_{11}$$

Then use Approach 2 for $A_{:,1}B_{11}$.

Remark (Analysis). Here is the analysis of page faults for this approach:

- $A_{:,1}$: 256 columns, $1024 \times 256 / 65536 = 4$ pages.
- $A_{:,1}, \dots, A_{:,4}$: $4 \times 4 = 16$ page faults in calculating $C_{:,1}$

- For A : 16×4 page faults
- B : 16 page faults
- C : 16 page faults

3.2.2 LAPACK

BLAS defines only operations such as matrix-matrix products. LAPACK –Linear Algebra PACKage, based on BLAS, which can solve

- Systems of linear equations
- Least-squares solutions of linear systems of equations
- Eigenvalue problems
- Singular value problems

Subroutines in LAPACK classified as three levels:

1. **Driver routines:** Subroutines that solve complete problems, such as linear systems, eigenvalue problems, and singular value problems.
2. **Computational routines:** Subroutines that perform a specific computation such as LU factorization, QR factorization, and Hessenberg reduction.
3. **Auxiliary routines:** Subroutines that perform auxiliary computations, which is kind of extension of BLAS.

Definition 3.2.3 (Name of LAPACK Driver/Computational Routines). Name of LAPACK Driver and Computational Routines has the form of

XYYZZZ

- X: data type (S: single, D: double, C: complex, Z: complex double)
- YY: the type of matrix
 - GB: general banded
 - GE: general (i.e. unsymmetric, in some cases rectangular)

Note. Banded matrix: a band of nonzeros along diagonals

$$\begin{pmatrix} * & * & 0 & 0 & 0 \\ * & * & * & 0 & 0 \\ 0 & * & * & * & 0 \\ 0 & 0 & * & * & * \\ 0 & 0 & 0 & * & * \end{pmatrix}$$

- ZZZ: the type of computation performed
 - SV: simple driver of solving general linear systems.
 - TRF: Factorization.
 - TRS: Use the factorization to solve $Ax = b$ by forward or backward substitution.

– CON: Estimate the reciprocal of the condition number.

Example. SGEV: Simple driver of solving single general linear systems.

Example. SGBSV: Simple driver of solving single general banded linear systems.

For the block algorithm in LAPACK, LU factorization DGETRF takes $O(n^3)$ operations, speed in megaflops (MFLOPS) is defined as

$$\text{MFLOPS} = \frac{\text{number of floating point operations}}{\text{time in seconds} \times 10^6}$$

	No. of CPUs	Block size	100	1000
Dec Alpha Miata	1	28	172	370
Compaq AlphaServer DS-20	1	28	353	440
IBM Power 3	1	32	278	551
IBM PowerPC	1	52	77	148
Intel Pentium II	1	40	132	250
Intel Pentium III	1	40	143	297
SGI Origin 2000	1	64	228	452
SGI Origin 2000	4	64	190	699
Sun Ultra 2	1	64	121	240
Sun Enterprise 450	1	64	163	334

Table 3.1: Performance of block LU factorization DGETRF in LAPACK on different machines.

We can see that block algorithms are not very effective for small-sized problems. In the Intel Pentium III, the CPU clock time is 550 MHz, so the theoretical peak performance is just almost achieved for $n = 1000$, which means that the block algorithm is very effective for large-sized problems.

Note (ATLAS). There is also some software packages for optimized BLAS, such as ATLAS (Automatically Tuned Linear Algebra Software), which can automatically tune the parameters of the block algorithm to achieve better performance. Which means compiled for your architecture, things related to your CPU, size of cache, RAM.